

NutriAI - Automated Diet Plan Builder

NutriAI is a full-stack, offline-runnable Streamlit application that generates a 7-day, 3-meal-a-day plan under the 60-second target while enforcing clinical condition filters, allergens, diet modes, diversity, and macro/micronutrient analytics.

Item	Implementation
Runtime	Streamlit UI + pure-Python planning package; runs with streamlit run code/app.py
Offline data	5,200 meal candidates in data/food_database.csv
Clinical rules	IBS, GERD, type 2 diabetes, hypertension, allergens, cross-contamination, diet and cultural constraints
Outputs	7-day plan, explain tables, RDA analytics, benchmarks, downloadable CSV, required persona checks

Architecture

Layer	Responsibility
Data layer	Loads offline food snapshot, RDA table, and clinical rule JSON. No network/API key required during grading.
Safety layer	Applies hard exclusions before ranking: diet compatibility, cultural constraints, exact allergen tags, cross-contamination risks, FODMAP, GERD, G
Retrieval layer	Embeds meal text and user profile with deterministic dense hashed vectors; scores candidates by cosine similarity.
Ranking layer	Optimizes calorie fit, micronutrient priorities, clinical fit, semantic similarity, and diversity penalties.
Analytics layer	Computes per-meal and per-day nutrients; compares against age/sex RDA/AI targets and flags gaps below 80%.
UI layer	Streamlit dashboard with profile form, plan table, nutrient views, explanations, benchmarks, persona tests, and CSV export.

Data Sources and Rule References

- USDA FoodData Central API guide for nutrition-field structure: <https://fdc.nal.usda.gov/api-guide>
- NIH Office of Dietary Supplements for RDA/AI definitions and nutrient references: <https://ods.od.nih.gov/HealthInformation/nutrientrecommendations/>
- NHLBI DASH eating plan for sodium cap and potassium/calcium/magnesium/fiber emphasis: <https://www.nhlbi.nih.gov/education/dash-eating-plan>
- Monash Low FODMAP program for IBS/FODMAP filtering concepts: <https://www.monashfodmap.com/>
- NIDDK celiac disease guidance for gluten cross-contact handling:
<https://www.niddk.nih.gov/health-information/digestive-diseases/celiac-disease/eating-diet-nutrition>

BAX-423 Technique Choices

Technique	Where used	Why it fits
Bloom filter sketching	Allergen and cross-contamination pre-screening in nutriai/bloom.py and dry_utils/membership.py	Fast, memory-efficient membership tests before exact safety checks. This demonstrates s
Dense embeddings	HashedFeature embeddings in nutriai/embedding.py	Transforms noisy meal text and user clinical/diet profile into vectors without needing m
Recommendation ranking	Multi-stage scorer in nutriai/planner.py	Balances calorie target, micronutrient priorities, clinical fit, semantic similarity, and dive

Benchmarks

Metric	Value
Food records scanned	5,200
Safe records for first persona	2,816
Bloom filter bits / hashes	95 bits / 6 hashes
Embedding dimensions	160
Candidate vectors scored	2,816
Total generation time	1.9985 seconds

The implementation intentionally avoids heavyweight infrastructure in the live demo path. The app still demonstrates big-data patterns by scanning a 5,200-row offline snapshot, sketching allergen membership, embedding candidate text, and ranking thousands of safe records per profile.

Required Persona Results

Persona	Time	Diversity	Clinical	Allergy	Diet	Nutrition	60 sec
Priya	2.00s	100.0	PASS	PASS	PASS	PASS	PASS
Ravi	3.09s	100.0	PASS	PASS	PASS	PASS	PASS
Mei	2.58s	100.0	PASS	PASS	PASS	PASS	PASS
James	2.29s	100.0	PASS	PASS	PASS	PASS	PASS

Persona-Specific Evidence

- Priya: vegetarian, lactose-free, low-FODMAP plan with no garlic/onion/wheat high-FODMAP meals and no dairy exposure.
- Ravi: gluten-free GERD plan that excludes high-acid tomato/citrus/fried/spicy triggers and pork.
- Mei: vegan, tree-nut-free, low-glycemic plan using fortified foods and low-GI ranking.
- James: pescatarian, soy-free, low-sodium/DASH-oriented plan with potassium, magnesium, and omega-3 priority.

Explainability

Every selected meal includes a ranking explanation with calorie fit, nutrient fit, clinical fit, similarity, and diversity factors. Every excluded example lists the exact reason, such as allergen detected, cross-contamination risk, high-FODMAP food for IBS, GERD trigger, high-GI food for diabetes, or high-sodium food for hypertension.

Limitations and Demo Script

Limitations

- The offline nutrition snapshot is deterministic and USDA-style, but not a live FoodData Central pull. This keeps grading runnable without API keys.
- Rules are conservative class-project approximations and should be reviewed by a clinician or registered dietitian before real-world use.
- The hashed embedding model is FAISS-ready in spirit but uses numpy cosine retrieval to avoid native binary install risk during demo.
- Meal serving multipliers are used for calorie targeting; production software would use ingredient-level recipe scaling and inventory constraints.

10-Minute Demo Flow

- 1 Open the app with streamlit run code/app.py and show the profile sidebar.
- 2 Load Priya and generate the plan; point out zero lactose, low-FODMAP exclusions, and 21 unique meals.
- 3 Open Nutrients and show per-day calories plus RDA comparison.
- 4 Open Explain and show selected-meal ranking reasons and excluded-food reasons.
- 5 Open Benchmarks and identify Bloom filter sketching, embeddings, and ranking timings.
- 6 Run Required Persona Tests and show all PASS statuses.

Submission Run Commands

Task	Command
Install	<code>pip install -r code/requirements.txt</code>
Run app	<code>streamlit run code/app.py</code>
Run tests	<code>python -m unittest discover -s code/tests -v</code>

Educational project only. Not medical advice.